Mechanism representation: sensitivity to input budget

Top-10 (%) with fixed K2 knowledge, 216 candidates; same native observation within each budget protocol.

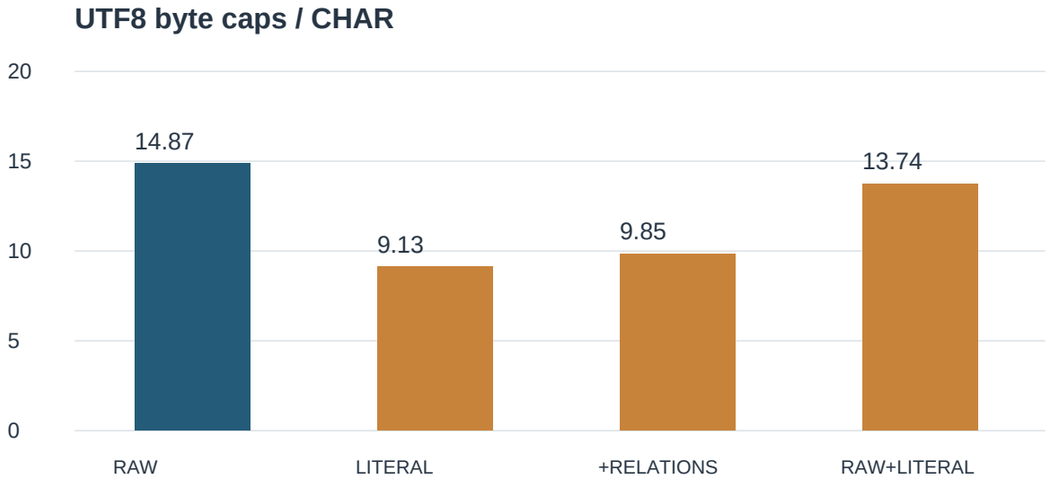

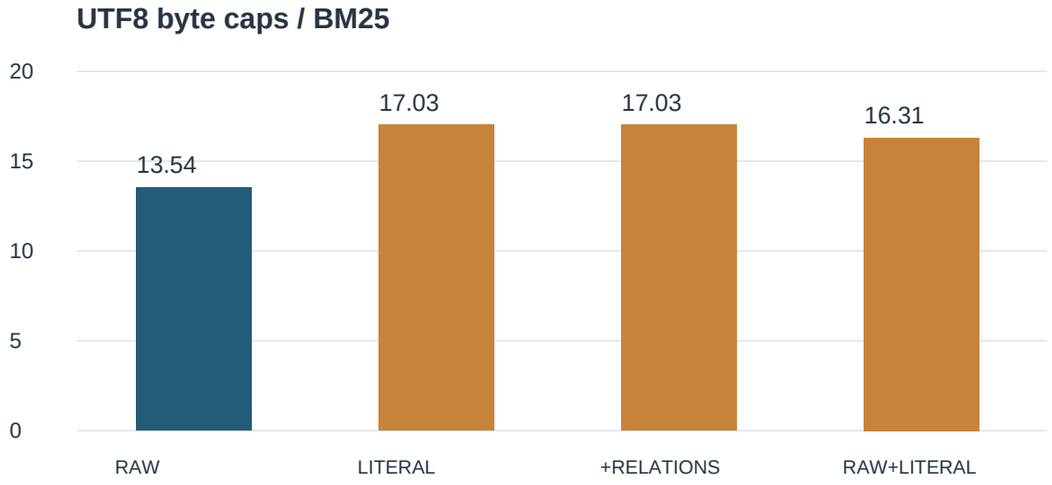

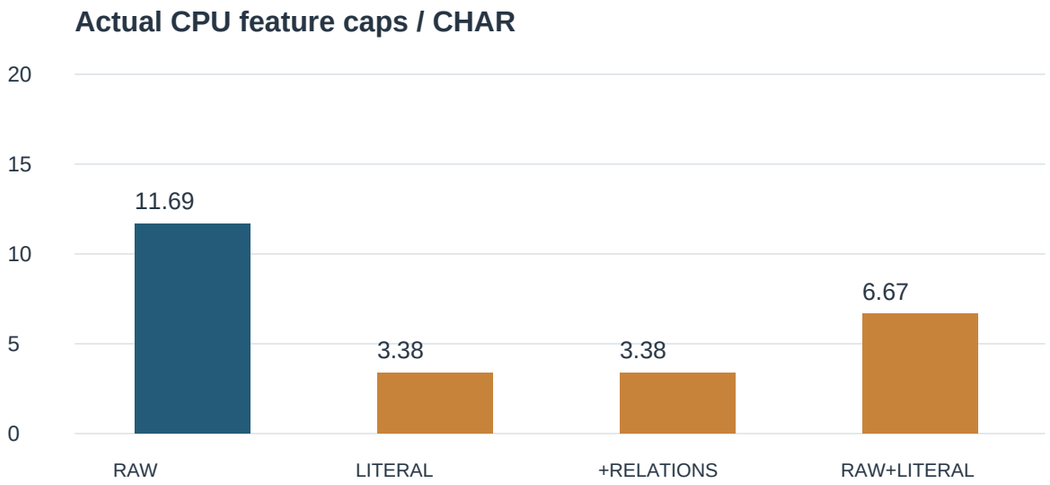

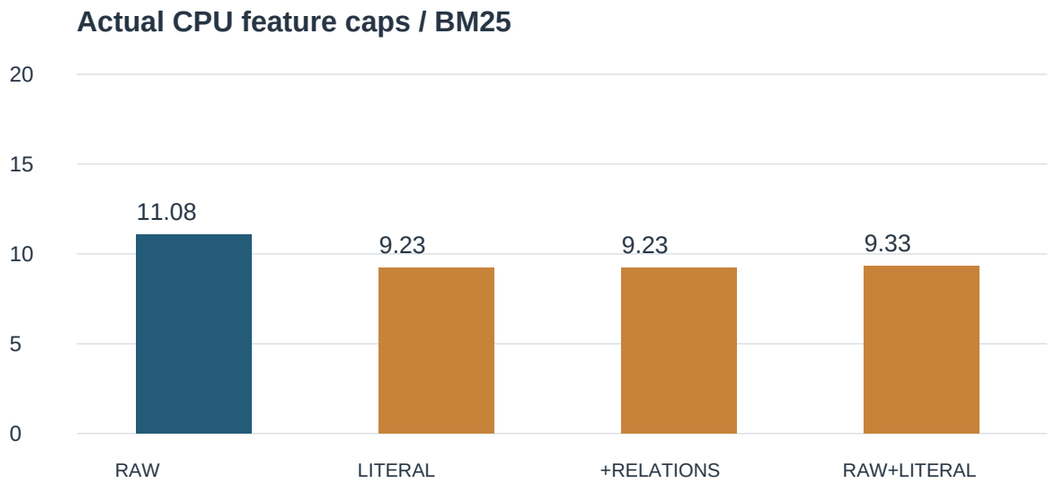

Feature units: character 3/4/5-gram occurrences and BM25 word tokens; these are NOT GTE tokens.

Relation markers have no BM25 knowledge postings. A zero delta does not test semantic relation reasoning.

Deterministic CPU runs; exposed development data only. No new neural representation scores are plotted.