Actual CPU projection training curves: seed42, ONE view

Each line averages 7 overlapping training folds. Loss definitions differ; levels do not establish cross-method superiority.

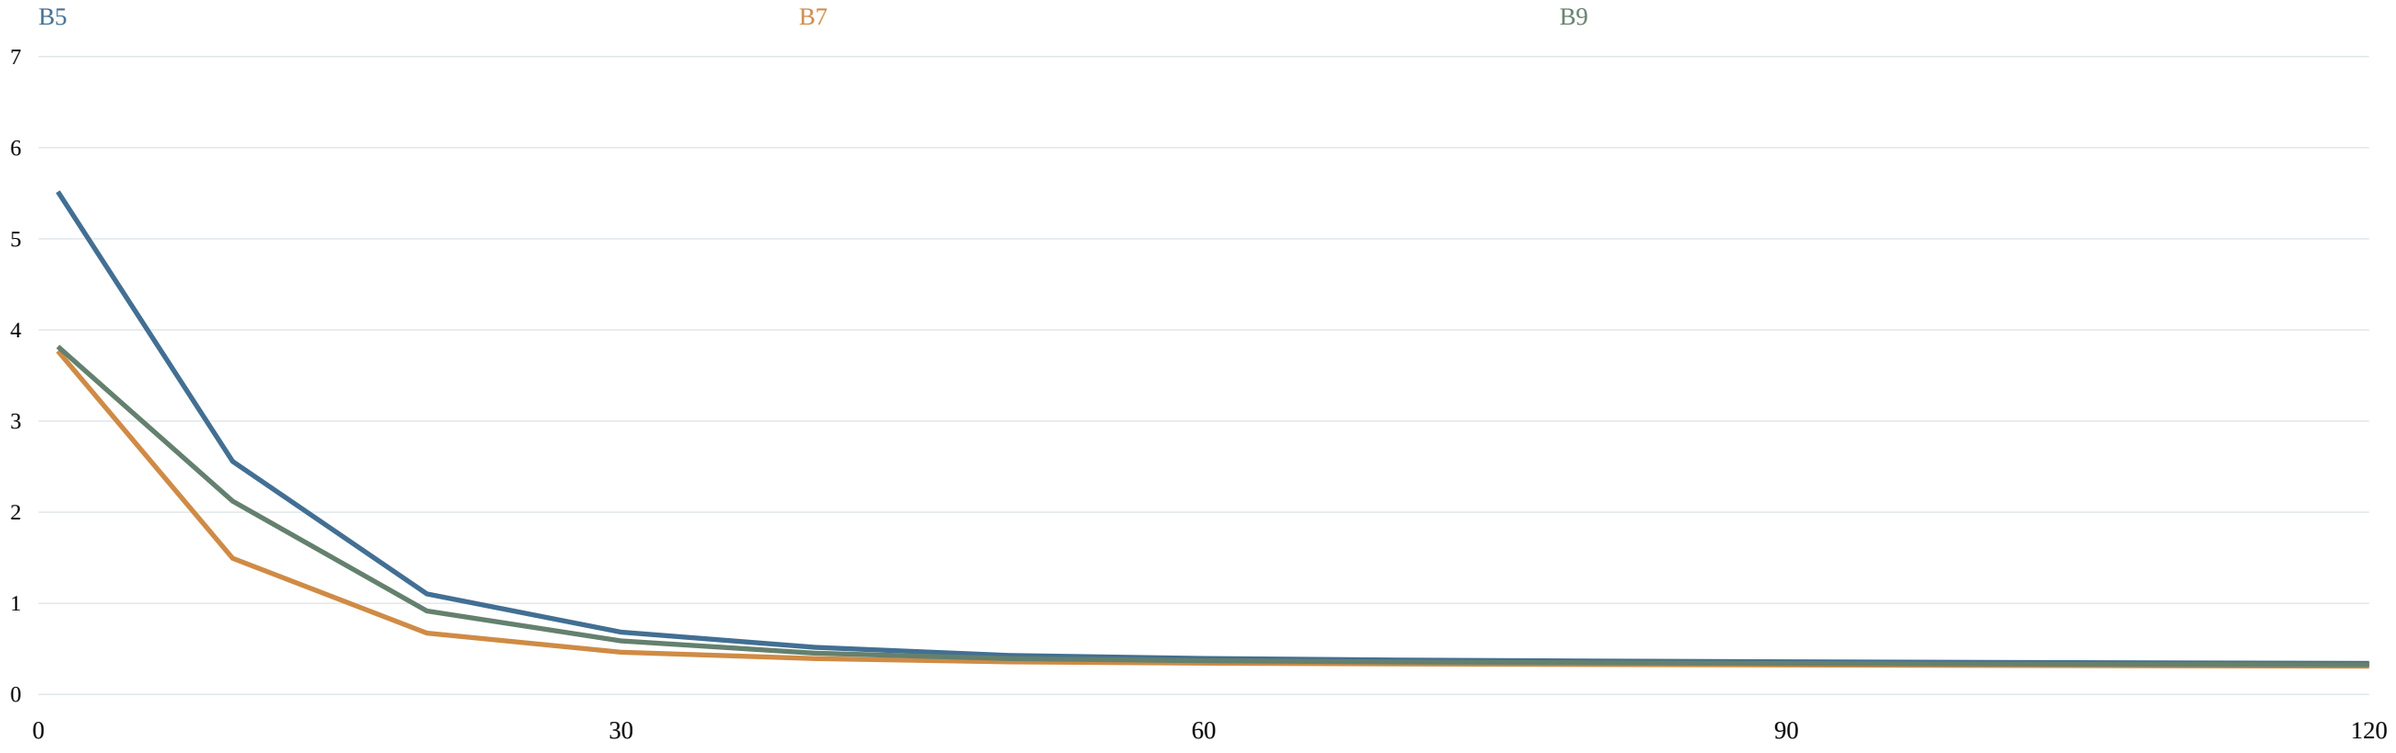

120 fixed Adam updates, no early stopping or heldout selection. Full logs/checkpoints saved; optimization convergence is not independent generalization.