The 32 lost hits are restored by removing the seen supervision component

Diagnostic component removal using the original lambda/gamma and saved scores; these are not formal model results.

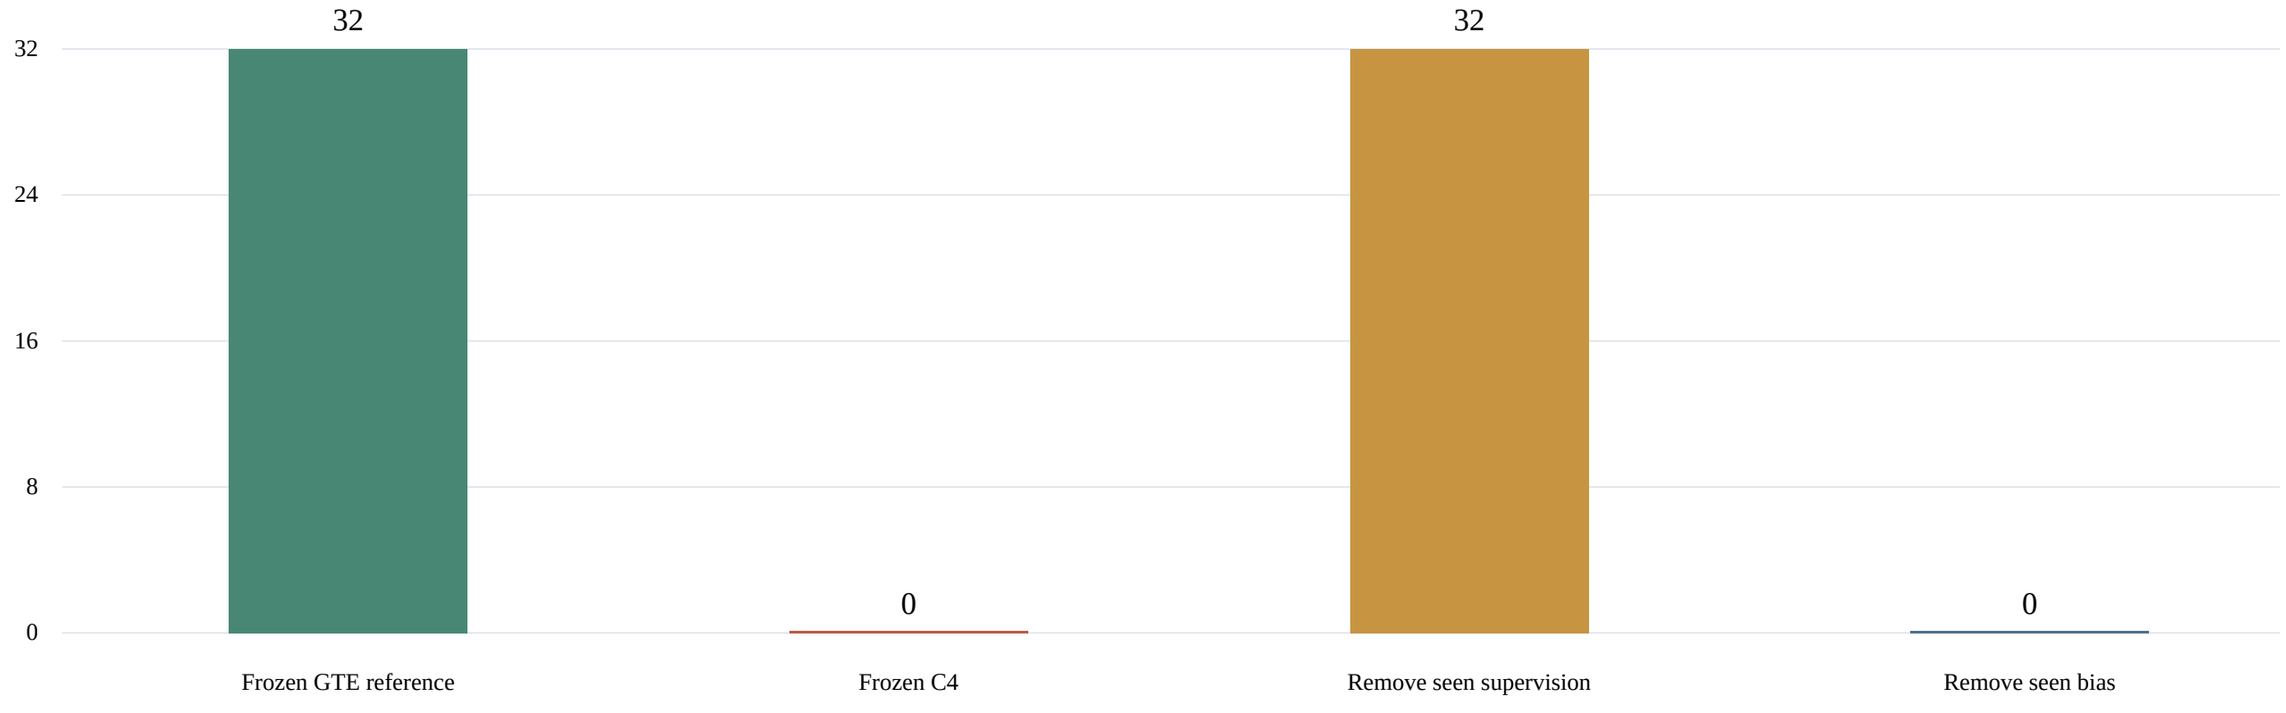

All 32 ONE losses have gamma=0; none uses the two-step validation fallback. No exact Gold-score ties.

This identifies a computational ranking cause, not an attack-semantic or annotation cause.