Supervised accuracy and unseen ranking remain in tension

ONE view; actual query-weighted metrics from saved 216-candidate predictions; bars are 3-seed means with seed SD.

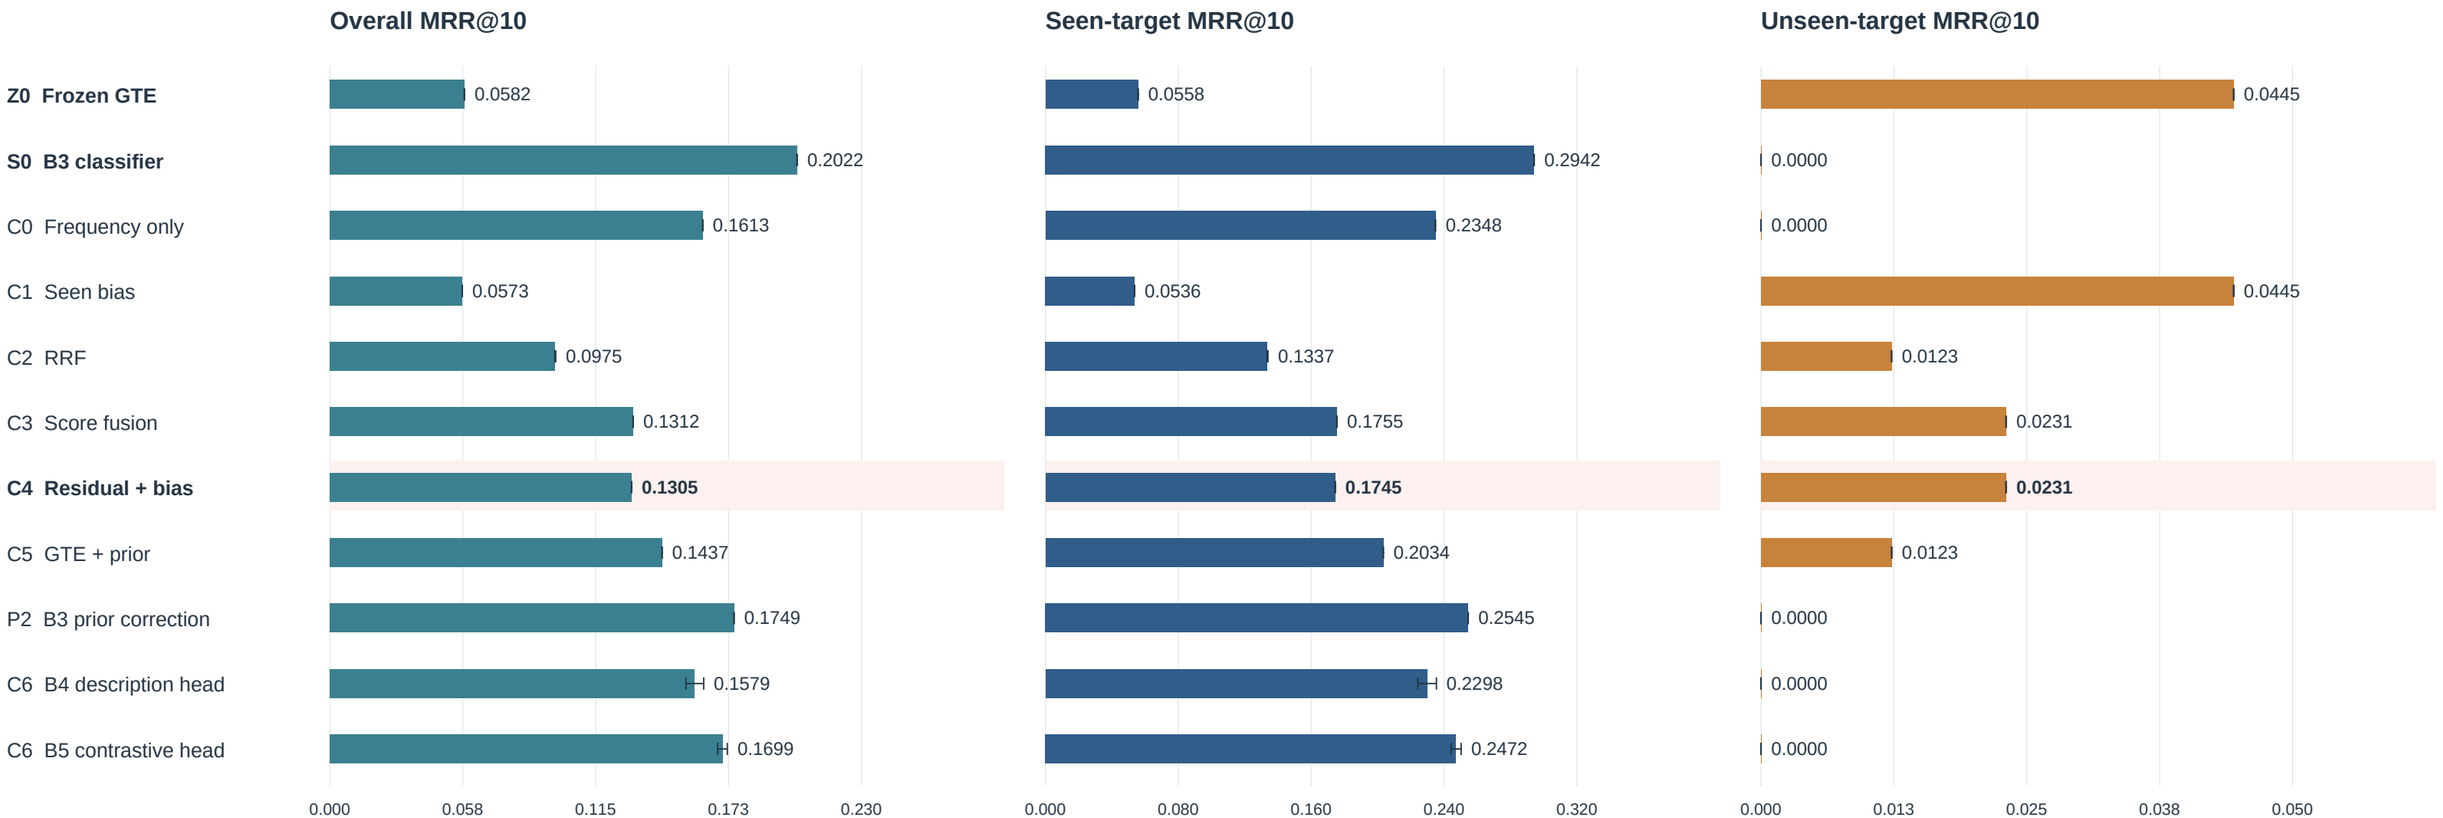

C4 restores unseen MRR above S0, retaining 59.3% of S0 seen MRR and 51.8% of Z0 unseen MRR.

Z0/C0/C1/C5 are deterministic repetitions for pairing; seed SD is not an independent-attack confidence interval.

7 exposed scenarios / 195 logical steps / 975 queries per condition. Seen targets: 670; unseen targets: 450, including 145 mixed queries.

Candidate rankings are never filtered by Gold. No new model training, GPU forward, or LLM call.